

Automated Model Building and Goodness-of-fit via Quantile Regression

Haim Bar

September 9, 2022

Abstract

This repository contains code and data used in the paper *Automated Model Building and Goodness-of-fit via Quantile Regression* by Bar, Booth, and Wells (in preparation). Given P predictors x_i and n observations for each x_i and a response variable y , the goal is to build a model, $y = f(x_1, \dots, x_P)$ where $f()$ consists of combinations of powers of the x_i 's, which fits the data well across multiple quantiles.

1 Prerequisites

In order to run the code you must first install the **QREM** package. Since **QREM** has a model selection option for cases in which the number of predictors is large you also need to install the packages **edgefinder** and **SEMMS**. The recommended way to install these packages is from the GitHub repository:

```
devtools::install_github("haimbar/edgefinder")
devtools::install_github("haimbar/SEMMS")
devtools::install_github("haimbar/QREM")
```

The model building algorithm is implemented in a function called *fitQRloop* in the file `runQREM.R`. The function takes five arguments:

- `M` the data matrix with P columns and n rows.
- `qns` The quantile which will be used in the fitting algorithm.
- `minDiff` The minimal improvement in the overall goodness of fit in order to accept a new term.
- `maxdeg` The maximum degree of any term in the model.
- `maxrows` The maximum number of possible terms up to degree `maxdeg`.

The file `initSim.R` contains the values we used by default. It also contains three other variables which are used by **QREM** in the fitting process:

- `mxm` The maximum number of segments in the partition of the selected variable.
- `alphaQ` The level of the goodness of fit test.
- `plotit` A Boolean variable which tells the function *flatQQplot* whether to show intermediate diagnostic plots for each accepted new term in the model.

```
qns <- 1:5/6
k <- length(qns)
minDiff <- 4
maxdeg <- 15
maxrows <- 5000
mxm <- 30
```

```
alphaQ <- 0.01
plotit <- FALSE
```

2 A Univariate Example

The file Code/Univariate02.R contains the code for example #1 in the paper, where

$$f(x) = x^5 e^{-x}$$

and the random noise is normally distributed with mean 0 and standard deviation which grows linearly with x , $sd = 0.25(x + 0.05)$.

```
source("Code/initSim.R") # set up some global parameters
qns <- round(qns, digits=3)
N <- 5000
set.seed(211111)
x <- runif(N, min=0, max=4*pi)
y <- exp(-x)*x^5 + rnorm(N, 0, 0.25*(x+0.05)) # EXAMPLE 1 in the paper
M <- data.frame(y=y, x1=x)
res <- fitQRloop(M=M, qn = qns, maxdeg = maxdeg, minDiff = minDiff) # ❶
pdf("Figures/Uni02.pdf", width=5, height=5)
plot(x, y, cex=0.5, pch=19, col="grey66", axes=F)
axis(1); axis(2); grid()
for (i in 1:k) {
  lines(sort(x), res$qremFit[[i]]$fitted.mod$fitted.values[order(x)],
        col=2)
}
dev.off()
```

The fitting algorithm is invoked in the line denoted by ❶ in the code above. The fitted quantile regression models are shown as red curves in Figure 1.

3 A Multivariate Example

The program in Code/multivariate04.R generates a data matrix with 2,000 observations and 4 predictors, but only x_1, x_2, x_4 are related to the response, according to:

$$y = x_1 x_2 x_4 + \epsilon$$

and the random errors are i.i.d. $\epsilon \sim N(0, 0.1^2)$

```
source("Code/initSim.R")
set.seed(211013)
N <- 2000
M <- matrix(runif(N*4, min = -1, max=3), nrow=N, ncol=4)
y <- M[,1]*M[,2]*M[,4] + rnorm(N, 0, 0.1)
M <- as.data.frame(cbind(y, M))
colnames(M) <- c("y", "x1", "x2", "x3", "x4")
res <- fitQRloop(M=M, qn = qns, maxdeg = maxdeg, minDiff = minDiff,
                 maxrows = maxrows)
pdf("Figures/multivariate04.pdf", width=5, height=5)
i <- which(qns == 0.5)
plot(M$y, res$qremFit[[i]]$fitted.mod$fitted.values, cex=0.6, pch=19,
      xlab="Observed", ylab="predicted", axes=FALSE, col="grey66")
abline(0, 1, col=3, lwd=2); axis(1); axis(2)
```




Figure 1: The fitted QR model for $q = 1/6, \dots, 5/6$ where the true model is $f(x) = x^5 e^{-x}$ with random errors i.i.d. $N(0, [0.25(x + 0.05)]^2)$.

The fitted model found by our algorithm is:

$$x_1 * x_2 * x_4 + I(x_4^2) + x_2 * x_4 + I(x_2^2) + x_1 * x_4 + x_1 * x_2 + I(x_1^2) + x_4 + x_2 + x_1$$

The predicted values are very close to the observed one, as can be seen in Figure 2.

Examples which involve more complicated models with more than one predictor may take a few minutes to run, and sometimes even longer.




Figure 2: Fitted QR model vs. observed values for $q = 1/2$ where the true model is $f(x_1, x_2, x_3, x_4) = x_1 x_2 x_4$ with random errors i.i.d. $N(0, 0.1^2)$.

4 Case Studies

The repository contains several case studies:

- Concrete.R - predicting the strength of concrete. (Note that fitting a model to this dataset is time-consuming. Saved results can be found in concreteResultsR2s.RData).
- mpgreg.R - which factors contribute to gasoline consumption. (We run it 10 times, and it takes a minute or two to finish). The file MPG.py contains a python program which uses TensorFlow to fit the MPG data. The code was obtained from the TensorFlow documentation.
- bankNotes.R (a classification example, see <https://archive.ics.uci.edu/ml/datasets/banknote+authentication>).
- Lidar.R - Lidar readings data (univariate, demonstrating data augmentation).
- uscrime.R - FBI rape rate data by state (demonstrating regression discontinuity).
- Ozone.R - Ozone data.

In this section we use the MPG data

The fitted model is:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	20.677	0.1335	154.911	2.42e-257
I(x6 ²)	0.509	0.0296	17.207	9.09e-45
x5	-0.446	0.0538	-8.285	6.38e-15
x6	1.937	0.0885	21.892	7.55e-61
x3	-2.690	0.1504	-17.883	3.94e-47
x2	-0.423	0.1403	-3.018	2.80e-03
x1(4.5,6.5]	-0.714	0.1217	-5.865	1.36e-08
x1(6.5,8.5]	-1.227	0.2670	-4.594	6.78e-06
x4	-6.166	0.1504	-41.005	2.92e-115
x5:x6	0.699	0.0548	12.746	3.28e-29
x6:x3	-1.083	0.1228	-8.817	1.76e-16
x6:x2	-0.092	0.1124	-0.819	4.14e-01
x6:x1(4.5,6.5]	0.548	0.1125	4.870	1.95e-06
x6:x1(6.5,8.5]	1.732	0.2360	7.339	2.78e-12
x1(4.5,6.5]:x4	3.933	0.1792	21.947	4.92e-61
x1(6.5,8.5]:x4	5.483	0.2101	26.101	1.78e-74

The predicted values are plotted versus the observed one in Figure 3.

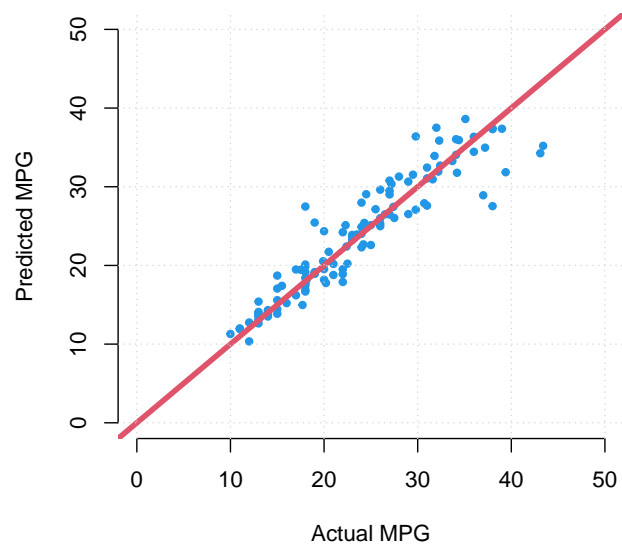


Figure 3: Fitted QR model vs. observed values for $q = 1/2$ for the MPG data.